

Hello this is @Ranjeet_Kumbhar, Enjoy the Notebook

GitHub:https://github.com/RanjeetKumbhar01/TE_IT_ML_ASSIGNMENTS_SPPU

Question

1. Data preparation: Download heart dataset from following link.
<https://www.kaggle.com/zhaoyingzhu/heartcsv> Perform following operation on given dataset.
 - a) Find Shape of Data
 - b) Find Missing Values
 - c) Find data type of each column
 - d) Finding out Zero's
 - e) Find Mean age of patients

f) Now extract only Age, Sex, ChestPain, RestBP, Chol. Randomly divide dataset in training (75%) and testing (25%). Through the diagnosis test I predicted 100 report as COVID positive, but only 45 of those were actually positive. Total 50 people in my sample were actually COVID positive. I have total 500 samples. Create confusion matrix based on above data and find I. Accuracy

II. Precision

III. Recall

IV. F-1 score

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import seaborn as sns

pwd
'/kaggle/working'

df = pd.read_csv('../input/heartcsv/Heart.csv')

df.head()

    Unnamed: 0  Age  Sex      ChestPain  RestBP  Chol  Fbs  RestECG
MaxHR \
0           1   63     1        typical     145    233    1       2
150
```

1	2	67	1	asymptomatic	160	286	0	2
108	3	67	1	asymptomatic	120	229	0	2
129	4	37	1	nonanginal	130	250	0	0
187	5	41	0	nontypical	130	204	0	2
172								
	ExAng	Oldpeak	Slope	Ca	Thal	AHD		
0	0	2.3	3	0.0	fixed	No		
1	1	1.5	2	3.0	normal	Yes		
2	1	2.6	2	2.0	reversable	Yes		
3	0	3.5	3	0.0	normal	No		
4	0	1.4	1	0.0	normal	No		

a) Find Shape of Data

```
df.shape #303, 15
```

```
(303, 15)
```

b) Find Missing Values

```
df.isnull().sum()
```

```
Unnamed: 0      0
Age            0
Sex            0
ChestPain     0
RestBP         0
Chol           0
Fbs            0
RestECG        0
MaxHR          0
ExAng          0
Oldpeak        0
Slope          0
Ca             4
Thal           2
AHD            0
dtype: int64
```

```
df.count()
```

```
Unnamed: 0    303
Age          303
```

```
Sex          303
ChestPain    303
RestBP       303
Chol         303
Fbs          303
RestECG      303
MaxHR        303
ExAng        303
Oldpeak      303
Slope         303
Ca            299
Thal          301
AHD           303
dtype: int64
```

c) Find data type of each column

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 15 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Unnamed: 0   303 non-null    int64  
 1   Age          303 non-null    int64  
 2   Sex          303 non-null    int64  
 3   ChestPain    303 non-null    object  
 4   RestBP       303 non-null    int64  
 5   Chol         303 non-null    int64  
 6   Fbs          303 non-null    int64  
 7   RestECG      303 non-null    int64  
 8   MaxHR        303 non-null    int64  
 9   ExAng        303 non-null    int64  
 10  Oldpeak      303 non-null    float64 
 11  Slope        303 non-null    int64  
 12  Ca            299 non-null    float64 
 13  Thal          301 non-null    object  
 14  AHD           303 non-null    object  
dtypes: float64(2), int64(10), object(3)
memory usage: 35.6+ KB
```

```
df.dtypes

Unnamed: 0      int64
Age             int64
Sex             int64
ChestPain      object
```

```

RestBP          int64
Chol           int64
Fbs            int64
RestECG        int64
MaxHR          int64
ExAng          int64
Oldpeak       float64
Slope          int64
Ca             float64
Thal           object
AHD            object
dtype: object

```

d) Finding out Zero's

`df==0`

	Unnamed: 0	Age	Sex	ChestPain	RestBP	Chol	Fbs
RestECG	\						
0	False	False	False	False	False	False	False
1	False	False	False	False	False	False	True
2	False	False	False	False	False	False	True
3	True	False	False	False	False	False	True
4	False	False	True	False	False	False	True
..
298	True	False	False	False	False	False	True
299	True	False	False	False	False	False	False
300	True	False	False	False	False	False	True
301	False	False	True	False	False	False	True
302	True	False	False	False	False	False	True
	MaxHR	ExAng	Oldpeak	Slope	Ca	Thal	AHD
0	False	True	False	False	True	False	False
1	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False
3	False	True	False	False	True	False	False
4	False	True	False	False	True	False	False

```

...
298 False True False False True False False
299 False True False False False False False
300 False False False False False False False
301 False True True False False False False
302 False True True False False False False

[303 rows x 15 columns]

df[df==0]

      Unnamed: 0  Age  Sex ChestPain  RestBP  Chol  Fbs  RestECG  MaxHR
ExAng \
0       NaN  NaN  NaN  NaN  NaN  NaN  NaN  NaN  NaN
0.0
1       NaN  NaN  NaN  NaN  NaN  NaN  0.0  NaN  NaN
NaN
2       NaN  NaN  NaN  NaN  NaN  NaN  0.0  NaN  NaN
NaN
3       NaN  NaN  NaN  NaN  NaN  NaN  0.0  0.0  NaN
0.0
4       NaN  NaN  0.0  NaN  NaN  NaN  0.0  NaN  NaN
0.0
...
...
298     NaN  NaN  NaN  NaN  NaN  NaN  0.0  0.0  NaN
0.0
299     NaN  NaN  NaN  NaN  NaN  NaN  NaN  0.0  NaN
0.0
300     NaN  NaN  NaN  NaN  NaN  NaN  0.0  0.0  NaN
NaN
301     NaN  NaN  0.0  NaN  NaN  NaN  0.0  NaN  NaN
0.0
302     NaN  NaN  NaN  NaN  NaN  NaN  0.0  0.0  NaN
0.0

      Oldpeak  Slope  Ca  Thal  AHD
0       NaN  NaN  0.0  NaN  NaN
1       NaN  NaN  NaN  NaN  NaN
2       NaN  NaN  NaN  NaN  NaN
3       NaN  NaN  0.0  NaN  NaN
4       NaN  NaN  0.0  NaN  NaN
...
...
298     NaN  NaN  0.0  NaN  NaN
299     NaN  NaN  NaN  NaN  NaN
300     NaN  NaN  NaN  NaN  NaN
301    0.0  NaN  NaN  NaN  NaN
302    0.0  NaN  NaN  NaN  NaN

[303 rows x 15 columns]

```

```
(df == 0).sum()

Unnamed: 0      0
Age            0
Sex           97
ChestPain     0
RestBP         0
Chol          0
Fbs           258
RestECG       151
MaxHR         0
ExAng         204
Oldpeak       99
Slope          0
Ca            176
Thal          0
AHD           0
dtype: int64
```

e) Find Mean age of patients

```
np.mean(df['Age'])
54.43894389438944

df.Age.mean()
54.43894389438944
```

f) Now extract only Age, Sex, ChestPain, RestBP, Chol. Randomly divide dataset in training (75%) and testing (25%).

```
df.columns
Index(['Unnamed: 0', 'Age', 'Sex', 'ChestPain', 'RestBP', 'Chol',
       'Fbs',
       'RestECG', 'MaxHR', 'ExAng', 'Oldpeak', 'Slope', 'Ca', 'Thal',
       'AHD'],
      dtype='object')

data = df[['Age', 'Sex', 'ChestPain', 'RestBP', 'Chol']]
#Cross validation
```

```
train,test = train_test_split(data,test_size=0.25,random_state=1)

train.shape
(227, 5)

test.shape
(76, 5)
```

Through the diagnosis test I predicted 100 report as COVID positive, but only 45 of those were actually positive. Total 50 people in my sample were actually COVID positive. I have total 500 samples. Create confusion matrix based on above data and find I. Accuracy

II. Precision

III. Recall

IV. F-1 score

type(predicted)

`numpy.ndarray`

Confusion Matrix

		Real Label	
		Positive	Negative
Predicted Label	Positive	True Positive (TP)	False Positive (FP)
	Negative	False Negative (FN)	True Negative (TN)

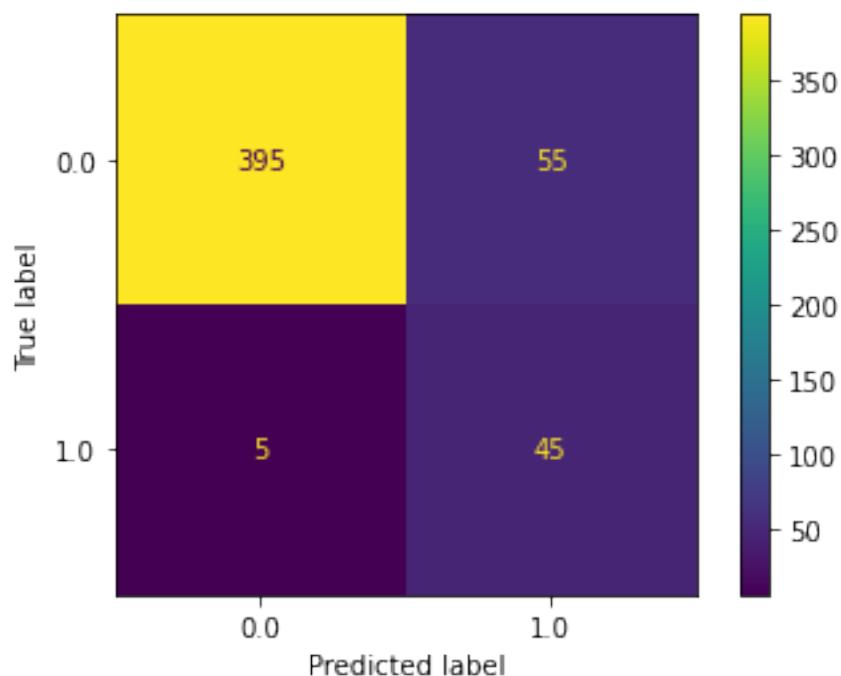
Precision = $\frac{\sum \text{TP}}{\sum \text{TP} + \text{FP}}$

↓

Recall = $\frac{\sum \text{TP}}{\sum \text{TP} + \text{FN}}$

Accuracy = $\frac{\sum \text{TP} + \text{TN}}{\sum \text{TP} + \text{FP} + \text{FN} + \text{TN}}$

```
from sklearn.metrics import ConfusionMatrixDisplay
ConfusionMatrixDisplay.from_predictions(actual,predicted)
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x7f133ed13d10>
```



```
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score
```

```
print(classification_report(actual,predicted))
```

	precision	recall	f1-score	support
0.0	0.99	0.88	0.93	450
1.0	0.45	0.90	0.60	50
accuracy			0.88	500
macro avg	0.72	0.89	0.76	500
weighted avg	0.93	0.88	0.90	500

```
accuracy_score(actual,predicted)
```

```
0.88
```