

Hello this is @Ranjeet_Kumbhar, Enjoy the Notebook

GitHub: https://github.com/RanjeetKumbhar01/TE_IT_ML_ASSIGNMENTS_SPPU Assignment on Regression technique Download temperature data from below link.

<https://www.kaggle.com/venky73/temperaturesof-india?select=temperatures.csv> This data consists of temperatures of INDIA averaging the temperatures of all places month wise. Temperatures values are recorded in CELSIUS

a. Apply Linear Regression using suitable library function and predict the Month-wise temperature.

b. Assess the performance of regression models using MSE, MAE and R-Square metrics

c. Visualize simple regression model.

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt
# import matplotlib.pyplot as plt

data = pd.read_csv("../input/temperatures-of-india/temperatures.csv")
df = data

data.describe()
```

	YEAR	JAN	FEB	MAR	APR	\
count	117.000000	117.000000	117.000000	117.000000	117.000000	
mean	1959.000000	23.687436	25.597863	29.085983	31.975812	
std	33.919021	0.834588	1.150757	1.068451	0.889478	
min	1901.000000	22.000000	22.830000	26.680000	30.010000	
25%	1930.000000	23.100000	24.780000	28.370000	31.460000	
50%	1959.000000	23.680000	25.480000	29.040000	31.950000	
75%	1988.000000	24.180000	26.310000	29.610000	32.420000	
max	2017.000000	26.940000	29.720000	32.620000	35.380000	
	MAY	JUN	JUL	AUG	SEP	
OCT \						
count	117.000000	117.000000	117.000000	117.000000	117.000000	
117.000000						
mean	33.565299	32.774274	31.035897	30.507692	30.486752	
29.766581						
std	0.724905	0.633132	0.468818	0.476312	0.544295	
0.705492						
min	31.930000	31.100000	29.760000	29.310000	29.070000	
27.900000						

25%	33.110000	32.340000	30.740000	30.180000	30.120000
29.380000					
50%	33.510000	32.730000	31.000000	30.540000	30.520000
29.780000					
75%	34.030000	33.180000	31.330000	30.760000	30.810000
30.170000					
max	35.840000	34.480000	32.760000	31.840000	32.220000
32.290000					

	NOV	DEC	ANNUAL	JAN-FEB	MAR-MAY
JUN-SEP \					
count	117.000000	117.000000	117.000000	117.000000	117.000000
	117.000000				
mean	27.285470	24.608291	29.181368	24.629573	31.517607
	31.198205				
std	0.714518	0.782644	0.555555	0.911239	0.740585
	0.420508				
min	25.700000	23.020000	28.110000	22.250000	29.920000
	30.240000				
25%	26.790000	24.040000	28.760000	24.110000	31.040000
	30.920000				
50%	27.300000	24.660000	29.090000	24.530000	31.470000
	31.190000				
75%	27.720000	25.110000	29.470000	25.150000	31.890000
	31.400000				
max	30.110000	28.010000	31.630000	28.330000	34.570000
	32.410000				

	OCT-DEC
count	117.000000
mean	27.208120
std	0.672003
min	25.740000
25%	26.700000
50%	27.210000
75%	27.610000
max	30.030000

```
data.head()
```

```
4 1905 22.00 22.83 26.68 30.01 33.32 33.25 31.44 30.68 30.12  
30.67
```

	NOV	DEC	ANNUAL	JAN-FEB	MAR-MAY	JUN-SEP	OCT-DEC
0	27.31	24.49	28.96	23.27	31.46	31.27	27.25
1	26.31	24.04	29.22	25.75	31.76	31.09	26.49
2	26.08	23.65	28.47	24.24	30.71	30.92	26.26
3	26.36	23.63	28.49	23.62	30.95	30.66	26.40
4	27.52	23.82	28.30	22.25	30.00	31.33	26.57

```
data.tail()
```

	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG
SEP	\								
112	2013	24.56	26.59	30.62	32.66	34.46	32.44	31.07	30.76
	31.04								
113	2014	23.83	25.97	28.95	32.74	33.77	34.15	31.85	31.32
	30.68								
114	2015	24.58	26.89	29.07	31.87	34.09	32.48	31.88	31.52
	31.55								
115	2016	26.94	29.72	32.62	35.38	35.72	34.03	31.64	31.79
	31.66								
116	2017	26.45	29.46	31.60	34.95	35.84	33.82	31.88	31.72
	32.22								

	OCT	NOV	DEC	ANNUAL	JAN-FEB	MAR-MAY	JUN-SEP	OCT-DEC
112	30.27	27.83	25.37	29.81	25.58	32.58	31.33	27.83
113	30.29	28.05	25.08	29.72	24.90	31.82	32.00	27.81
114	31.04	28.10	25.67	29.90	25.74	31.68	31.87	28.27
115	31.98	30.11	28.01	31.63	28.33	34.57	32.28	30.03
116	32.29	29.60	27.18	31.42	27.95	34.13	32.41	29.69

```
type(data)
```

```
pandas.core.frame.DataFrame
```

```
data.shape
```

```
(117, 18)
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 117 entries, 0 to 116  
Data columns (total 18 columns):  
 #   Column    Non-Null Count  Dtype    
---  --    
 0   YEAR      117 non-null    int64   
 1   JAN       117 non-null    float64  
 2   FEB       117 non-null    float64  
 3   MAR       117 non-null    float64
```

```
4 APR      117 non-null   float64
5 MAY      117 non-null   float64
6 JUN      117 non-null   float64
7 JUL      117 non-null   float64
8 AUG      117 non-null   float64
9 SEP      117 non-null   float64
10 OCT     117 non-null   float64
11 NOV     117 non-null   float64
12 DEC     117 non-null   float64
13 ANNUAL  117 non-null   float64
14 JAN-FEB 117 non-null   float64
15 MAR-MAY 117 non-null   float64
16 JUN-SEP 117 non-null   float64
17 OCT-DEC 117 non-null   float64
dtypes: float64(17), int64(1)
memory usage: 16.6 KB
```

```
count = (data["JAN"]==22).sum()
print(count)
```

```
1
```

```
column = data
count = column[column == 0].count()
print(count)
```

```
YEAR      0
JAN       0
FEB       0
MAR       0
APR       0
MAY       0
JUN       0
JUL       0
AUG       0
SEP       0
OCT       0
NOV       0
DEC       0
ANNUAL    0
JAN-FEB   0
MAR-MAY   0
JUN-SEP   0
OCT-DEC   0
dtype: int64
```

```
data.isnull().sum()
```

```
YEAR      0
JAN       0
FEB       0
```

```
MAR      0
APR      0
MAY      0
JUN      0
JUL      0
AUG      0
SEP      0
OCT      0
NOV      0
DEC      0
ANNUAL   0
JAN-FEB  0
MAR-MAY  0
JUN-SEP  0
OCT-DEC  0
dtype: int64
```

```
data.isnull().head()
```

```
    YEAR    JAN    FEB    MAR    APR    MAY    JUN    JUL    AUG
SEP \
0  False  False  False  False  False  False  False  False  False
1  False  False  False  False  False  False  False  False  False
2  False  False  False  False  False  False  False  False  False
3  False  False  False  False  False  False  False  False  False
4  False  False  False  False  False  False  False  False  False
      OCT    NOV    DEC  ANNUAL  JAN-FEB  MAR-MAY  JUN-SEP  OCT-DEC
0  False  False  False  False  False  False  False  False
1  False  False  False  False  False  False  False  False
2  False  False  False  False  False  False  False  False
3  False  False  False  False  False  False  False  False
4  False  False  False  False  False  False  False  False
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 117 entries, 0 to 116
Data columns (total 18 columns):
 #   Column    Non-Null Count  Dtype  
 --- 
 0   YEAR      117 non-null    int64  
 1   JAN       117 non-null    float64
 2   FEB       117 non-null    float64
 3   MAR       117 non-null    float64
```

```

4    APR      117 non-null   float64
5    MAY      117 non-null   float64
6    JUN      117 non-null   float64
7    JUL      117 non-null   float64
8    AUG      117 non-null   float64
9    SEP      117 non-null   float64
10   OCT      117 non-null   float64
11   NOV      117 non-null   float64
12   DEC      117 non-null   float64
13  ANNUAL     117 non-null   float64
14 JAN-FEB     117 non-null   float64
15 MAR-MAY     117 non-null   float64
16 JUN-SEP     117 non-null   float64
17 OCT-DEC     117 non-null   float64
dtypes: float64(17), int64(1)
memory usage: 16.6 KB

data.head()

    YEAR    JAN    FEB    MAR    APR    MAY    JUN    JUL    AUG    SEP
OCT \
0  1901  22.40  24.14  29.07  31.91  33.41  33.18  31.21  30.39  30.47
29.97
1  1902  24.93  26.58  29.77  31.78  33.73  32.91  30.92  30.73  29.80
29.12
2  1903  23.44  25.03  27.83  31.39  32.91  33.00  31.34  29.98  29.85
29.04
3  1904  22.50  24.73  28.21  32.02  32.64  32.07  30.36  30.09  30.04
29.20
4  1905  22.00  22.83  26.68  30.01  33.32  33.25  31.44  30.68  30.12
30.67


    NOV    DEC  ANNUAL  JAN-FEB  MAR-MAY  JUN-SEP  OCT-DEC
0  27.31  24.49  28.96   23.27   31.46   31.27   27.25
1  26.31  24.04  29.22   25.75   31.76   31.09   26.49
2  26.08  23.65  28.47   24.24   30.71   30.92   26.26
3  26.36  23.63  28.49   23.62   30.95   30.66   26.40
4  27.52  23.82  28.30   22.25   30.00   31.33   26.57

# x = data.iloc[:,1:6]
# y = data.iloc[:, -1:]
x = data["YEAR"]
y = data["ANNUAL"]


plt.plot(x,y, 'o')

[<matplotlib.lines.Line2D at 0x7fe03a5342d0>]

```



```
sns.scatterplot(x=x,y=y,data=df)  
<AxesSubplot:xlabel='YEAR' , ylabel='ANNUAL'>
```



```
type(x)  
pandas.core.series.Series
```

```

x.shape
(117,)
x = x.values
x = x.reshape(117,1)
x.shape
(117, 1)
type(x)
numpy.ndarray
x_train, x_test, y_train, y_test = train_test_split(x,
y,test_size=0.25)

print(f"x Training dataset: {x_train.shape}")
print(f"y Training dataset: {y_train.shape}")
print(f"x test dataset: {x_test.shape}")
print(f"y test dataset: {y_test.shape}")

x Training dataset: (87, 1)
y Training dataset: (87,)
x test dataset: (30, 1)
y test dataset: (30,)

model = LinearRegression()
model.fit(x_train,y_train)
LinearRegression()

```


$$y = w \cdot x + b$$

```


model.coef_ #w
array([0.01181101])
model.intercept_ #b
6.004942319521042
y_pred = model.predict(x_test)
y_pred.shape
(30,)

```

```
plt.scatter(x_train, y_train, color='blue')
plt.plot(x_test, y_pred, color='red', linewidth=3)
plt.title("Temperature vs Year")
plt.xlabel("Year")
plt.ylabel("Temperature")
plt.show()
```



```
sns.regplot(data=df,x=x_train,y=y_train,)
<AxesSubplot:ylabel='ANNUAL'>
```



```
from sklearn.metrics import  
mean_absolute_error, mean_squared_error, r2_score  
print(f"MSE: {mean_squared_error(y_test,y_pred)}")  
print(f"MAE: {mean_absolute_error(y_test,y_pred)}")  
print(f"R-Sqaure : {r2_score(y_test,y_pred)}")  
  
MSE: 0.25534763392126797  
MAE: 0.2911243760198463  
R-Sqaure : 0.5321685895690408
```